

Fig. 8a: Per-Class Severity Generalisation Gap
(Train: 0.007" + 0.014", Test: 0.021")




Fig. 8b: SVM vs XGBoost Noise Robustness
(Crossover Threshold Detection)

